

SAMYAK KAMBLE

Mumbai, MH • +91 78758 02373 • samyakkamble3016@gmail.com • [LinkedIn](#) • [GitHub](#) • [Portfolio](#)

PROFESSIONAL SUMMARY

Data Engineering Intern with hands-on experience building production-style ETL pipelines, ML forecasting systems, and cloud analytics architectures using Python, SQL, Apache Spark, Databricks, and Azure. Developed an end-to-end Retail Analytics ETL + ML Forecasting Pipeline processing 500K+ records with Prophet/LSTM forecasting, churn prediction, FastAPI serving, and Streamlit dashboards. Currently at Amdox Technologies working on PostgreSQL, Power BI, MLflow, Airflow, and Docker-based data engineering workflows.

TECHNICAL SKILLS

Programming & Querying: Python, SQL, Pandas, NumPy

Data Engineering: ETL Pipelines, Data Modeling, Data Warehousing, Data Cleaning, Feature Engineering, Batch Processing

Big Data & Cloud: Apache Spark, Databricks, Microsoft Azure, PySpark, Apache SQL, Hadoop, AWS Basics, Data Modelling

Databases: PostgreSQL, MySQL, SQL Server, MongoDB

Visualization & Tools: Power BI, Tableau, Streamlit, FastAPI, Docker, Airflow, MLflow, Git, GitHub, Jupyter, Matplotlib

EXPERIENCE

Data Engineering & Analytics Intern | Amdox Technologies

April 2026 – Present

- Built end-to-end ETL data pipelines for retail analytics — ingesting, cleaning, and transforming raw transaction data using Python, Pandas, SQL, and PostgreSQL.
- Architected production-style data workflow: ETL layer → PostgreSQL → FastAPI scoring API → Power BI / Streamlit dashboards → MLflow / Airflow MLOps orchestration.
- Implemented data quality checks, RFM metrics, and KPI generation; applied Apache Spark and Databricks for scalable batch processing on Azure cloud infrastructure.
- Containerized analytics stack with Docker; maintained structured GitHub repositories following production-grade versioning and documentation standards.

PROJECTS

Retail Analytics ETL + ML Forecasting Pipeline

Jan 2026 – Apr 2026

Python, SQL, Pandas, PostgreSQL, Prophet, LSTM, FastAPI, Streamlit, MLflow, Airflow, Docker, GitHub

- Built a production-style end-to-end data pipeline processing 500K+ retail transaction records across ingestion, cleaning, transformation, feature engineering, ML forecasting, API serving, and dashboard visualization.
- Engineered 20+ business and time-series features — TotalPrice, daily sales aggregation, RFM metrics, and customer behavior indicators — producing forecasting-ready datasets for downstream ML models.
- Developed and benchmarked Linear Regression, Prophet, and LSTM models for 30-day demand forecasting; built an RFM-based churn prediction classifier evaluated with precision, recall, and F1-score.
- Integrated MLflow for end-to-end experiment tracking and model reproducibility; exposed forecast outputs through FastAPI endpoints delivering real-time JSON predictions.
- Designed a Streamlit dashboard for revenue KPIs, sales trends, customer segments, and model forecasts; orchestrated ETL scheduling with Airflow DAGs and containerized the full stack using Docker.

Customer Shopping Behaviour Analysis

Sep 2025 – Dec 2025

SQL, Python, Pandas, Power BI

- Analyzed 50,000+ customer transactions using SQL and Pandas to surface spending patterns, seasonal trends, and discount impact; built an interactive Power BI dashboard for business KPI reporting and data storytelling.
- Built an interactive Power BI dashboard visualizing customer segments, sales KPIs, and discount effectiveness for business stakeholder reporting.

EDUCATION

BSc in Data Science | Mumbai University, Maharashtra

June 2024 – May 2027

Relevant Coursework: Database Management, Statistical Analysis, Data Structures, Machine Learning, Cloud Computing

CERTIFICATIONS

HackerRank SQL (Advanced) Certificate

April 2026

HackerRank Python (Advanced) Certificate

April 2026

Deloitte Data Analytics Job Simulation – Forage

July 2025

Tata Data Visualisation: Empowering Business with Effective Insights – Forage

June 2025

Artificial Intelligence Certificate – Sheth N.K.T.T. College

Mar 2026